

Exercise 0 – Notebook & Data Setup

Create a **new Databricks notebook** with default language **Python**.

1. Attach the notebook to a running cluster.
2. Load a reasonably **large dataset** into a DataFrame `df` (at least a few million records if possible).
3. Perform some basic checks:
 - o Show schema.
 - o Show a few rows.
 - o Count the records.
4. Create one or two **derived DataFrames** (e.g. filtered, aggregated) that you'll reuse later for caching experiments.

Deliverables for this exercise:

- A DataFrame `df` loaded from some data source.
 - At least one derived DataFrame, e.g. `df_filtered`, `df_agg`.
-

Exercise 1 – Warm-up: Actions vs Transformations

1. For `df`, define a transformation pipeline, e.g.:
 - o `filter` on some column.
 - o `withColumn` to add or transform a column.
 - o `select` to keep only a subset of columns.
2. **Do not** call any actions yet. Confirm there's no job triggered in the UI.
3. Call an **action** (e.g. `count()`) on the final DataFrame and observe:
 - o A job runs for the first time.
4. Call the same action again on the same DataFrame **without caching**:
 - o Observe that Spark recomputes the lineage (you will see another job).

Deliverables:

- Code for transformation pipeline.
 - Evidence (from the UI or logs) that the job ran twice.
-

Exercise 2 – Basic `cache()` with Performance Comparison

Take a **CPU-heavy / shuffle-heavy** transformation (like a `groupBy + agg`).

1. Measure the **time** taken to run a particular action (e.g., `count()`, `collect()` or `show()`):
 - o First **without caching**.
2. Now **cache** the resulting DataFrame.
3. Run the same action again and measure the time:

- Compare first and second runs.
- 4. Check if the DataFrame is cached using Spark's catalog APIs.
- 5. Explore:
 - How many actions benefit from the cache.
 - What happens if you **re-cache** or **unpersist**.

Deliverables:

- Timings (even approximate) for uncached vs cached action.
 - Code and comments summarizing your observations.
-

Exercise 3 – `persist()` with Different Storage Levels

1. Create a DataFrame `df_heavy` that's **expensive to compute** (e.g., multi-step pipeline with joins, aggregations, etc.).
2. Persist `df_heavy` with:
 - `MEMORY_ONLY`
 - `MEMORY_AND_DISK`
 - `DISK_ONLY`
3. For each storage level:
 - Persist the DataFrame.
 - Trigger an action to materialize it.
 - Measure the runtime of:
 - First action (materialization).
 - Second action (served from cache).
4. Check `df_heavy.storageLevel` after each persist.
5. If possible, observe memory usage / storage info in the Spark UI (Storage tab):
 - How much of the DataFrame fits in memory?
6. Unpersist the DataFrame and confirm it's removed from cache.

Deliverables:

- Code snippets showing different storage levels.
 - Short notes on performance and storage level behavior.
-

Exercise 4 – Caching Temporary Views & SQL Queries

1. Register a **temporary view** from your DataFrame, e.g.
`df.createOrReplaceTempView("my_table")`.
2. Run a **SQL query** using `%sql` on this view that is somewhat expensive (join, `groupBy`, etc.).
3. **Cache the table** using Spark SQL or PySpark:
 - e.g., `CACHE TABLE` or corresponding PySpark API.
4. Re-run the same SQL query:
 - Compare the runtime with and without cache.
5. Check:

- Whether the table is listed as cached in Spark's catalog.
6. Drop or uncache the table and confirm it is removed.

Deliverables:

- SQL and Python code related to caching the table.
 - Observations about performance benefits.
-

Exercise 5 – Cache Invalidation, Reuse, and Best Practices

Cache a DataFrame `df_cached`.

1. Perform multiple actions on `df_cached` and note performance.
2. Modify the **underlying data** (simulate this however is easiest in your environment):
 - Example: re-read the source file but with an extra filter, limit, or sample to mimic data changes.
3. Observe:
 - Does `df_cached` automatically reflect underlying data changes?
4. Unpersist `df_cached` and recreate it from the new data:
 - Cache again and rerun actions.
5. Think about and write down answers to:
 - When should you cache?
 - When should you avoid caching?
 - How do you avoid over-caching and memory pressure?

Deliverables:

- Code demonstrating cache reuse and unpersist.
-